

# Predicting Residential Home Sales Prices Using Regression Analysis

**Statistics 511**

**Prepared for: Professor Millie Mao  
Spring 2021**

**Rumil Legaspi <[rlegaspi20@apu.edu](mailto:rlegaspi20@apu.edu)>  
Nancy Huerta <[nhuerta20@apu.edu](mailto:nhuerta20@apu.edu)>**

## **Table of Contents**

<b>Purpose</b>	<b>3</b>
<b>Our Data</b>	<b>3</b>
Background on Dataset & Variables	3
<b>Part 1 - Model Estimation and Interpretation</b>	<b>4</b>
1a. Fitting a regression model estimating sales price using beds, baths, and garage size as predicting variables	4
1b. Interpretation of Coefficients	4
1c. Interpretation of Adjusted R-Squared = 0.54	5
<b>Part 2 - Prediction</b>	<b>5</b>
2a. Predicting the house sales price for a house with 3 bedrooms, 3 bathrooms, and a 2-car garage	5
2b. Calculating the 95% confidence interval	5
2c. Calculating the 95% prediction interval	5
<b>Part 3 - Hypothesis Testing</b>	<b>6</b>
3a. Checking the significance for each individual partial slope (independent variable)	6
3b. Conducting an F-test to check overall model significance	7
3c. Conducting Partial F tests	7
<b>Part 4 - Multicollinearity</b>	<b>8</b>
4a. Creating scatterplots and correlation matrices	9
4b. Removing Two Strongly Correlated Variables	11

## Purpose

We are conducting a multiple linear regression from the Real Estate Sales (APPENC07) dataset to analyze the relationship of the given features, bedrooms, bathrooms, and garage size, with the outcome variable, house sales price in a Midwestern city.

## Our Data

### Background on Dataset & Variables

Our dataset consists of *522 total transactions* from home sales during the year 2002.

Response Variable (Y)	Explanatory Variable 1 ( $X_1$ )	Explanatory Variable 2 ( $X_2$ )	Explanatory Variable 3 ( $X_3$ )
“house_price” sales price of residence (in dollars)	“beds” Number of bedrooms	“baths” Number of bathrooms	“garage_size” Number of cars the garage can hold

```
#Setting up our work environment
setwd("C:/Users/RUMIL/Desktop/APU/STAT 511 - Millie Mao (Applied Regression Analysis)/Project 2")
library(nortest)
library(olsrr)
library(car)
library(lmtest)
library(MASS)
library(tidyverse)
library(ggcorrplot)

#Loading in the text data
raw_data = read.table(file = "APPENC07.txt", header = FALSE, sep = "")

#Converting into tibble data frame for easier data analysis
house_data <- as_tibble(raw_data)
```

```
#Defining and renaming our Explanatory(X) and Response(Y) variables
house_data <- house_data %>% select(house_price = V2,
                                   beds = V4,
                                   baths = V5,
                                   garage_size = V7)

#Setting explanatory and response variables
house_price <- house_data %>% select(house_price) #Y
beds <- house_data %>% select(beds) #X1
baths <- house_data %>% select(baths) #X2
garage_size <- house_data %>% select(garage_size) #X3
```

## **Part 1 - Model Estimation and Interpretation**

### **1a. Fitting a regression model estimating sales price using beds, baths, and garage size as predicting variables**

```
#Using the lm function to fit a multiple regression model
house_lm <- lm(house_price ~ beds + baths + garage_size, data = house_data)

#Regression summary
summary(house_lm)

##
## Call:
## lm(formula = house_price ~ beds + baths + garage_size, data = house_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -249973  -55441  -16444   33862  423872
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -45886.3    17261.6  -2.658   0.0081 **
## beds         935.4      4966.4    0.188   0.8507
## baths       67818.9     5150.4   13.168  <2e-16 ***
## garage_size  67332.3     7176.3    9.383  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 93300 on 518 degrees of freedom
## Multiple R-squared:  0.5451, Adjusted R-squared:  0.5424
## F-statistic: 206.9 on 3 and 518 DF,  p-value: < 2.2e-16
```

### **1b. Interpretation of Coefficients**

#### ***Intercept & Partial Slopes***

From summarizing our multiple linear regression model we can see:

<b>Intercept</b>	<b>Bedrooms</b>	<b>Bathrooms</b>	<b>Garage Size</b>
$\beta_0$	$\beta_1$	$\beta_2$	$\beta_3$
-45886.3	935.4	67818.9	67332.3

and the estimated regression equation to be:

$$\hat{Y} = -45886.3 + 935.4X + 67818.9X + 67332.3X$$

The partial slopes in our summary indicate that when any one of the partial slopes **Increase by 1 unit**

and **other explanatory variables held constant and unchanged** we can expect:

- While holding our other explanatory variables Bathrooms and Garage Size constant and unchanged, when **Bedrooms** increase by 1 unit, we can expect our **house sales price** to increase by **roughly \$935.40**
- While holding our other explanatory variables Bedrooms and Garage Size constant and unchanged, when **Bathrooms** increases by 1 unit, we can expect our **house sales price** to increase by **roughly \$67,818.90**.
- While holding our other explanatory variables Bedrooms and Bathrooms constant and unchanged, when **Garage size** increases by 1 unit, we can expect our **house sales price** to increase by **roughly \$67,332.30**.

### 1c. Interpretation of Adjusted R-Squared = 0.54

Adjusted R squared value, similar to the R square value, tells us how much of the variability in our model is explained by our predictor variables (**Bedrooms, Bathrooms, Garage Size**), while also penalizing redundant or otherwise useless predictor variables, helping us to resist urges of adding too many variables into our model.

In this case our adjusted  $R^2$  of 0.54 tells us that about **54% of the variation in our response variable is explained by our 3 explanatory variables**.

## Part 2 - Prediction

### 2a. Predicting the house sales price for a house with 3 bedrooms, 3 bathrooms, and a 2-car garage

```
#We create an artificial observation where a given house has  
#3 Bedrooms, 3 Bathrooms, and a 2 car garage  
new_house_data <- data.frame(beds = 3, baths = 3, garage_size = 2)
```

### 2b. Calculating the 95% confidence interval

```
#confidence interval  
c1_house <- predict(house_lm, new_house_data, interval = "confidence", level = 0.95)  
c1_house
```

```
##          fit          lwr          upr  
## 1 295041.2 284025.7 306056.6
```

### Interpretation

This 95% confidence interval, when Bedrooms = 3, Bathrooms = 3, and Garage Size = 2, is from **\$284,025 to \$306,056**

When Bedrooms = 3, Bathrooms = 3, and Garage Size = 2, with 95% confidence we can expect our confidence interval to capture the average of house sales prices (response variable).

## 2c. Calculating the 95% prediction interval

```
#prediction interval
p1_house <- predict(house_lm, new_house_data, interval = "prediction", level = 0.95)
p1_house

##          fit          lwr          upr
## 1 295041.2 111422.3 478660
```

### Interpretation

From the results we can predict with 95% confidence that when there are 3 bedrooms, 3 bathrooms, and a garage that can hold 2 cars, the predicted house sales price will be **\$295,041** and fall somewhere between **\$111,422 to \$478,660**.

## Part 3 - Hypothesis Testing

### 3a. Checking the significance for each individual partial slope (independent variable)

Using a significance level of  $\alpha = 0.05$

Null Hypothesis:  $H_0: \beta_j = 0$  (slopes are showing no change),  $X_j$  is **not** linearly associated with Y, therefore the partial slope is **not significant**.

Alternative Hypothesis:  $H_1: \beta_j \neq 0$  (slopes are showing change),  $X_j$  is linearly associated with Y, therefore the partial slope is **significant**.

Table 3: Table Representation of Hypothesis Testing

Bedrooms ( $X_1$ )	Bathrooms ( $X_2$ )	Garage Size ( $X_3$ )
$0.8507 > \alpha = 0.05$	$<2e-16 < \alpha = 0.05$	$<2e-16 < \alpha = 0.05$
Fail to reject $H_0$	Reject $H_0$	Reject $H_0$
Not Significant	Significant	Significant

#### Bedroom variable:

The p-value of Bedroom is 0.8507 and is greater than our  $\alpha$  (accepted error) of 0.05, so we fail to reject our NULL hypothesis and must conclude with our NULL hypothesis. Stating that our partial slope, Bedrooms, does not show overall significance in our model.

#### Bathroom & Garage Size variables:

On the other hand, because the p-value of Bathroom and Garage size are both  $<2e-16$  and are incredibly smaller than our (accepted error) of 0.05, so we reject our NULL hypothesis and conclude with our alternative hypothesis. Our alternative hypothesis states that our partial slopes, Bathroom and Garage Size, shows overall significance in our model.

### 3b. Conducting an F-test to check overall model significance

Using a significance level of  $\alpha = 0.05$

**Null Hypothesis:**  $H_0: \beta_1 = \beta_2 = \beta_3 = 0$  (No partial slopes are significant). Shows no change, therefore does **not show** overall model significance.

**Alternative Hypothesis:**  $H_1: \beta_1 = \beta_2 = \neq 0$  (At least one partial slope is significant). Shows change, therefore **showing** overall model significance.

```
#We can use the qt() to find our critical value and compare with our t-value (test statistic)
# We use 0.95 Because of our 95% confidence interval and 518 for our degrees of freedom
qt(0.975, 518)
```

```
## [1] 1.964554
```

```
#Checking for our f-value
anova(house_lm)
```

```
## Analysis of Variance Table
##
## Response: house_price
##          Df      Sum Sq   Mean Sq F value    Pr(>F)
## beds      1 1.6931e+12 1.6931e+12 194.515 < 2.2e-16 ***
## baths     1 2.9426e+12 2.9426e+12 338.057 < 2.2e-16 ***
## garage_size 1 7.6627e+11 7.6627e+11  88.032 < 2.2e-16 ***
## Residuals 518 4.5089e+12 8.7044e+09
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Table 4: Table Representation of F-test Hypothesis Testing

Test Statistic Type & Result	Bedrooms ( $X_1$ )	Bathrooms ( $X_2$ )	Garage Size ( $X_3$ )
F-value	194.515 > 1.96	338.057 > 1.96	88.032 > 1.96
P-value	2.2e-16 < 0.05	2.2e-16 < 0.05	2.2e-16 < 0.05
Result	Significant	Significant	Significant

Our p-value of < 2.2e-16 being less than our alpha and our F values being larger than our critical value tells us we can **reject** our NULL hypothesis and conclude with our alternative hypothesis, that **at least one** of our predictor variables **shows** overall model significance.

### 3c. Conducting Partial F tests

**Which variable is actually contributing?**

Conducting partial F tests is important to see if the number of bathrooms ( $X_2$ ) and garage size ( $X_3$ ) are jointly



significant.

Using a significance level of  $\alpha = 0.05$

**Null Hypothesis:  $H_0$ : There is no change** when adding certain predictors to the significance of our model

**Alternative Hypothesis:  $H_1$ : There is change** when adding certain predictors towards the significance of our model.

```
#full model
house_lm

##
## Call:
## lm(formula = house_price ~ beds + baths + garage_size, data = house_data)
##
## Coefficients:
## (Intercept)      beds      baths  garage_size
##   -45886.3      935.4    67818.9     67332.3

#reduced model without bathrooms and garage size
bed_lm <- lm(house_price ~ beds, data = house_data)
```

We now *compare* our reduced model with our complete model

```
anova(bed_lm, house_lm)

## Analysis of Variance Table
##
## Model 1: house_price ~ beds
## Model 2: house_price ~ beds + baths + garage_size
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      520 8.2178e+12
## 2      518 4.5089e+12  2 3.7089e+12 213.04 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Since the p-value is **2.2e-16** is less than our significance level of **0.05** we see that bathroom and garage size are both jointly significant and therefore we can reject the null hypothesis, indicating there is significance in keeping both bathroom and garage size in our model.

In effect, we are concluding that bathroom and garage size are predictors that do contribute information in the prediction of house sales price and therefore should be retained in the model.

## **Part 4 - Multicollinearity**

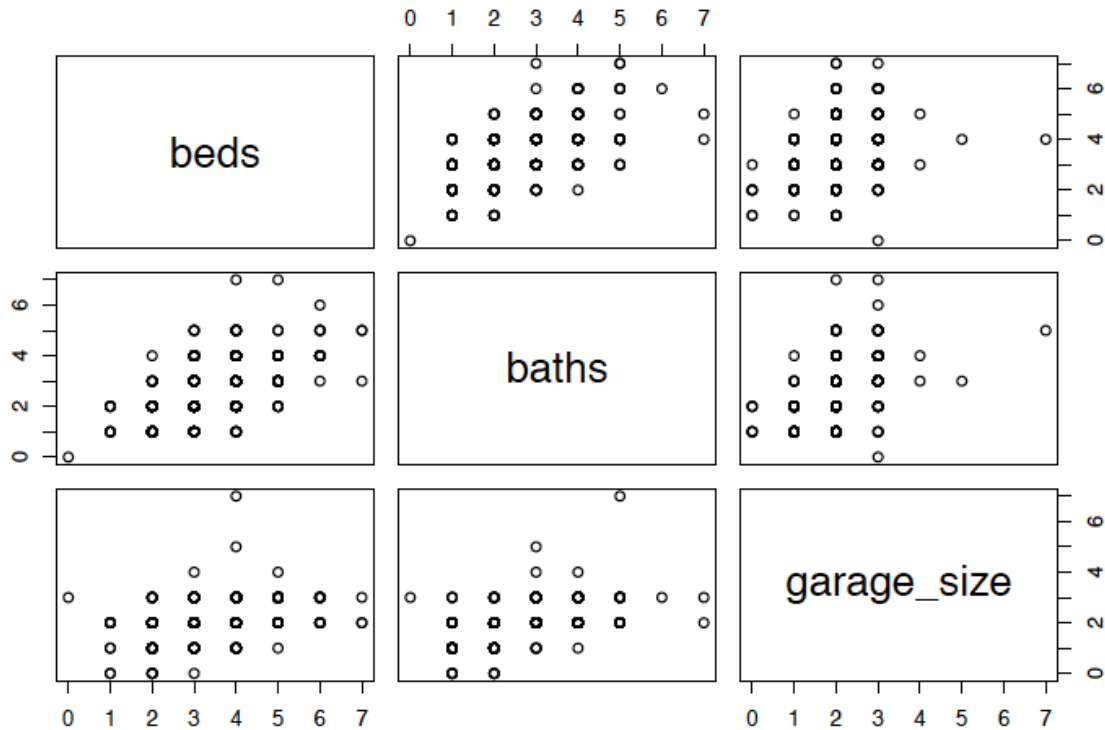
### **Why bother with multicollinearity?**

Having multicollinearity is problematic because by having multiple correlated predictor variables, it becomes harder for our model to attribute significance to our predictor variables. It creates redundant and duplicate information, thereby negatively affecting the results of our regression model.



#### 4a. Creating scatterplots and correlation matrices

```
#Plotting a scatterplot matrix **(why does it look symmetrical?)
scat_matrix <- c(beds, baths, garage_size) %>%
  data.frame() %>%
  plot()
```



```
scat_matrix
```

```
## NULL
```

```
#Correlation Matrix
corr_matrix <- c(beds, baths, garage_size) %>%
  data.frame() %>%
  cor()
```

```
corr_matrix
```


```
##           beds    baths garage_size
## beds      1.0000000 0.5834469  0.3168137
## baths      0.5834469 1.0000000  0.4898981
## garage_size 0.3168137 0.4898981  1.0000000
```

```
ggcorr_matrix <- ggcorrplot(corr_matrix, hc.order = TRUE, type = "lower", lab = TRUE,
  outline.col = "white",
  ggtheme = ggplot2::theme_gray,
  colors = c("#6D9EC1", "white", "#E46726"))
```

```
#Printing both matrices
scat_matrix
```

```
## NULL
```

```
ggcorr_matrix
```



```
summary(house_lm)
```

```
##
## Call:
## lm(formula = house_price ~ beds + baths + garage_size, data = house_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -249973  -55441  -16444   33862  423872
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
##
```

```
## (Intercept) -45886.3    17261.6   -2.658    0.0081 **
## beds        935.4      4966.4    0.188    0.8507
## baths       67818.9    5150.4   13.168   <2e-16 ***
## garage_size 67332.3    7176.3    9.383   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 93300 on 518 degrees of freedom
## Multiple R-squared:  0.5451, Adjusted R-squared:  0.5424
## F-statistic: 206.9 on 3 and 518 DF,  p-value: < 2.2e-16
```

```
corr_matrix
```

```
##           beds    baths garage_size
## beds      1.0000000 0.5834469  0.3168137
## baths      0.5834469 1.0000000  0.4898981
## garage_size 0.3168137 0.4898981  1.0000000
```

#### 4b. Removing Two Strongly Correlated Variables

A way to combat this is by removing a highly correlated predictor. From the correlation matrix and by looking at our correlation coefficient, we can see moderately positive relationship between bedrooms and bathrooms which might be worth further investigating.

We know there is some kind of multicollinearity issue with bed and baths, I am more interested in beds than baths. So I can remove baths from our model thereby correcting our multicollinearity issue.

```
#summary of original predictor
summary(house_lm)
```

```
##
## Call:
## lm(formula = house_price ~ beds + baths + garage_size, data = house_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -249973  -55441  -16444   33862  423872
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -45886.3    17261.6   -2.658    0.0081 **
## beds         935.4      4966.4    0.188    0.8507
## baths       67818.9    5150.4   13.168   <2e-16 ***
## garage_size 67332.3    7176.3    9.383   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 93300 on 518 degrees of freedom
## Multiple R-squared:  0.5451, Adjusted R-squared:  0.5424
## F-statistic: 206.9 on 3 and 518 DF,  p-value: < 2.2e-16
```

We notice bedrooms is not a significant variable from looking at the p-value, when in fact, the reality is **it should** be significant. Knowing this, we can check to see how well our model performs when removing bathroom since there is a multicollinearity issue.

```
#removing beds**
nobeds_lm <- lm(house_price ~ baths + garage_size, data = house_data)
summary(nobeds_lm)

##
## Call:
## lm(formula = house_price ~ baths + garage_size, data = house_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -249830  -55576  -15656   33933  423631
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -44091     14377   -3.067  0.00228 **
## baths           68321      4402  15.521 < 2e-16 ***
## garage_size    67391      7163   9.409 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 93210 on 519 degrees of freedom
## Multiple R-squared:  0.545, Adjusted R-squared:  0.5433
## F-statistic: 310.9 on 2 and 519 DF, p-value: < 2.2e-16
```

```
#removing baths
nobaths_lm <- lm(house_price ~ garage_size + beds, data = house_data)
summary(nobaths_lm)

##
## Call:
## lm(formula = house_price ~ garage_size + beds, data = house_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -352121  -66704  -28488   42621  529386
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -62861     19868   -3.164  0.00165 **
## garage_size    104753      7606  13.773 < 2e-16 ***
## beds           34804      4904   7.098 4.19e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 107700 on 519 degrees of freedom
## Multiple R-squared:  0.3928, Adjusted R-squared:  0.3904
## F-statistic: 167.9 on 2 and 519 DF, p-value: < 2.2e-16
```

In both cases we can see that by either removing bedroom or bathroom in our model, the predictors still remain significant but more importantly we can now see that bedrooms is in fact a significant predictor when removing the bathroom variable in our model concluding that we have addressed our issue of multicollinearity.